

Cloud Computing

MUCPD – DSIC – UPV

Virtualization

Overview

- **Introduction to virtualization**
- **History and evolution**
- **Types of virtualization**
 - **Hardware Virtualization**
 - **Software virtualization**
 - **OS-Level virtualization**
 - **Network virtualization**
- **Hypervisors: VM vs Containers**
- **Virtualization and the Cloud**
- **Pros and Cons**
- **Future trends**

Introduction

- **Definition**

Process of creating a software-based representation of something, rather than a physical one

- **Key Concepts**

- **Abstraction**


- Separation of resources and services from the physical delivery environment
 - Abstracts the hardware layer
 - Allows multiple virtual environments to run on a single physical system

- **Resource Sharing**

- All virtual environments share available physical resources
 - Potentially maximizes resource utilization

- **Isolation**

- Each virtual environment operates independently
 - A virtual environment cannot affect another virtual environment



Key Benefits

- **Improved Resource Utilization**
 - Sharing physical resource increases their utilization
 - From 15% utilization to 70-80% utilization
- **Cost Reduction**
 - Hardware reduction made possible, which reduces **CapEx** & **OpEx**
- **Simplified Management**
 - Virtualized resources are managed centrally
 - Easier to monitor performance, set updates, ...
- **Enhanced disaster recovery**
 - Virtual machines can be backed up fully
 - Recovery on different physical hardware
- **Environmental impact**

History and Evolution

- **1960s: IBM Mainframes and Time-Sharing**
 - E.g. CP-40, CP-67. Goal:
 - Maximize utilization of expensive hardware
- **1970s-1980s: Decline**
 - More accessible computing resources (PCs, Workstations).
 - Less focus on mainframes
- **1990s: Server Sprawl. Renewed interest**
 - Organizations accumulate servers (sprawl)
 - Virtualization to consolidate servers and increase efficiency
- **2000s: X86 virtualization**
 - VMWare & VirtualPC
 - Virtualization applied to widely used hardware → Game changer
- **2010-present: Cloud computing & Containers**
 - Orchestration tools,...

Early Virtualization-IBM Mainframes

- **CP-40 and CP-67 Projects**
 - Introduced the concept of Virtual Machines:
 - Multiple Operating Systems could run on a single physical mainframe
- **The idea of virtual machines**
 - A Machine within a Machine
 - Let users and their applications run in multiple ISOLATED environments
 - Sharing the physical resources
- **Time-sharing systems**
 - Allow multiple users simultaneous access to computing resources
 - Maximize utilization.
 - Improve service being delivered to more users
- **Key: ISOLATION**
 - In various degrees

Decline and Resurgence

- **1970s-1980s Decline**

- Rise of Personal Computers and Workstations
 - Computing power accessible to individuals with no need to access mainframes
 - Reduced need for shared resources

- **1990s: Resurgence**

- Multiple servers being deployed
 - Each dedicated to a specific task/app (why?)
 - Thus... underutilized hardware
- Inefficient hardware utilization
 - Increased cost to operate many physical servers
 - IT departments face increased costs: diversity of hardware and software
- Need for consolidation
 - Consolidation of servers through virtualization techniques:
 - simplify management.
 - Reduce costs

Virtualizing the X86 architecture

- **2000's breakthrough**

- Virtualization on standard hardware (The PC architecture)
 - Prior only on proprietary mainframes. Not useful for sprawl of servers
 - The PC architecture was (is) the basis for most PCs and servers
- Virtual PC (Microsoft)
 - Early fully software-emulated architecture
 - Very inefficient, as many resources employed to emulate a hardware machine
 - Very portable, as it was, essentially, a process interpreting machine code
- VMWare contribution
 - 1999: VMWare workstation. For workstations/PCs/Servers.
 - 2001: VMWare ESX Server. Type 1 Hypervisor for data centers.

- **Impact on Data Centers**

- Allowed virtualization of server environments (and consolidation)
- Groundwork for the development of cloud computing services
 - Efficiency needed to manage large amounts of VMs

Emergence of Cloud Computing

- **2006 Onwards**
 - AWS offered IT infrastructure services
- **Reliance on virtualization**
 - Virtualization allowed a scalable offering of on-demand computing resources as services
 - Key to allow providers to dynamically offer resources on demand
- **Shift in IT Paradigm**
 - Organizations move from OWNING physical devices to CONSUMING virtualized ones
 - Business can focus on their core
 - Outsource infrastructure management to third parties
 - Increased flexibility and reduced costs

Container Technologies

- **2013 Onwards**
 - Introduction of Docker
 - Simplified container management and deployment on a single node
- **Advantages over hardware virtualization**
 - Lightweight and fast start up
 - No boot time. Process launching
 - Improved application deployment
 - Images enabled consistency and portability
- **Impact on devops & microservices**
 - Facilitates CI/CD pipelines
 - Facilitates building microservice-based architectures for services.

Types of virtualization

- **Hardware Virtualization**
 - Creates virtual version of physical devices
 - Enables multiple OSs to run on a single physical machine within their own VM
- **Software Virtualization**
 - Focuses on virtualizing applications/software environments
 - E.g. Desktop virtualization
- **OS level virtualization**
 - Containers
 - Multiple isolated user-space instances running on a single OS kernel
- **Network Virtualization**
 - Abstracts physical network resources
 - Creates multiple virtual networks over a single physical network
 - Includes technologies like SDN (Software Defined Networks).

Hardware Virtualization

- **Full Virtualization**

- Hypervisor provides a complete simulation of the underlying hardware
- Guest OS runs unmodified.
 - Unaware it is running virtualized

- **Paravirtualization**

- Hypervisor does not fully simulate the hardware
 - Guest OS is modified (lightly) to communicate with hypervisor directly
 - Performance gains

- **Use Cases**

- Server consolidation
 - Multiple loads under same physical server
- Testing & Development Environments
 - Quickly spin-up/down multiple VMs for testing
 - Potentially on different environments

OS level virtualization

- **Containers**
 - Isolated user spaces
 - Share Host OS Kernel
- **Advantages**
 - Lightweight
 - Fast
 - Scalable
- **Technologies**
 - Docker
 - Podman
 - Containerd
 - LXC

Network Virtualization

- **Various elements**

- Virtual switches and routers
 - Implemented within the kernel
- Virtual interfaces
 - Implemented within the kernel
- VLANs
 - Part of the ethernet protocol
- Software Defined Networks
 - Separates control-plane from data-plane
 - Improves network programmability

- **Benefits**

- Improved network management
- Enhanced security
 - Sets up specific communication paths (0-trust networking on top)
- Flexibility

Hypervisors

A hypervisor, also known as a Virtual Machine Monitor (VMM), is software that creates and runs virtual machines by abstracting the underlying hardware resources

- **Type1 (Bare Metal)**
 - Runs directly on the host's hardware
 - ESXi, Hyper-V, XenServer,...
- **Type 2 (Hosted)**
 - On top of an OS
 - E.g. VMWare Workstation, VirtualBox, Parallels.
- **Responsibilities**
 - Resource allocation
 - Isolation
 - Management

Type 1 vs Type 2 Hypervisors

Type 1

- **Architecture:**
 - Install directly on the physical hardware.
 - hypervisor manages all hardware resources and provides virtual machines with direct access to physical hardware.
- **Performance:**
 - Offers better performance and efficiency since there's no underlying operating system.
- **Use Cases:**
 - Ideal for enterprise data centers, cloud providers, and environments where performance and scalability are critical.
- **Examples:**
 - VMware ESXi, Microsoft Hyper-V Server, Citrix XenServer.

Type 2

- **Architecture:**
 - Install on top of an existing operating system.
 - Relies on the host OS for device support and resource management.
- **Performance:**
 - May have lower performance due to additional overhead from the host OS.
- **Use Cases:**
 - Suitable for individual users, developers, and testing environments where ease of setup and flexibility are important.
- **Examples:**
 - VMware Workstation, Oracle VirtualBox, Parallels Desktop.

Types of Hypervisors

Characteristic	Type 1 (Bare-metal Hypervisor)	Type 2 (Hosted Hypervisor)
Architecture	Runs directly on hardware	Runs on a host OS
Performance	High performance	Lower performance
Use Cases	Data centers, cloud environments	Development, testing environments
Host OS Requirement	No	Yes
Example Hypervisors	VMware ESXi, Microsoft Hyper-V	VirtualBox, VMware Workstation

Virtual Machine

Software emulation of a Physical Computer

- **Components:**

- Virtual Hardware
 - Virtual CPU, memory storage
- Guest OS instance
 - Can be different from the host OS (for Type 2)
- Applications
 - On top of the guest OS

- **Properties**

- Isolation → Security, stability
- Flexibility
 - Different Guest OS
 - Legacy application support

VMs vs Containers

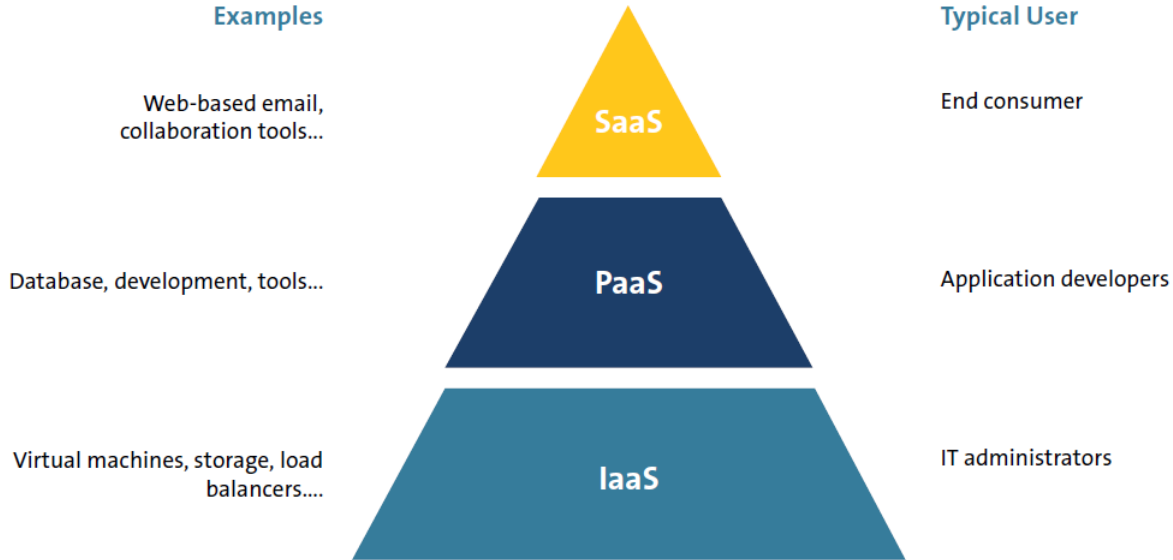
- **Structural:**
 - **VM:** Hypervisor-based, whole OS included
 - **Container:** Shares Kernel
- **Resource Usage**
 - **VM:** Higher overhead
 - **Container:** Lightweight/process-like overhead
- **Use Cases**
 - **VM:** Multi-OS environments. STRONG isolation
 - **Container:** Microservices, rapid deployment.

Virtualization: the backbone of the cloud

- **Abstraction of resources**
 - Managed by software, thus flexible.
- **Foundation for Cloud Services**
 - Much more difficult without it
- **Enables Scalability**
 - Rapid provisioning.
 - Allows Service providers to adapt to demand
- **Supports Multi-Tenancy**
 - Through Isolation

Cloud Service Models rehash


Figure 1: Cloud Pyramid²²



IaaS impact

- **All resources are virtualized**
- **No virtualization, no IaaS**
- **User controls OS & applications**
 - Full stack
- **Examples**
 - Amazon EC2
 - Microsoft Azure Virtual Machines
 - Google Compute Engine
 - OVH
 - Aire Networks
 - Private CPDs managed with OpenStack
 - ...

IaaS impact



PaaS impact

- **Provides Developer's support**
- **Abstracts OS management**
- **User controls applications**
- **Examples**
 - Heroku
 - Axebow/Kumori
 - Microsoft Azure App Service
 - Google App Engine
 - ...

SaaS impact

- **Delivers an application in the form of a service**
- **Abstracts managing application**
- **Application-specific user interactions**
 - User does not manage the application nor its deployment
- **Examples**
 - Salesforce
 - Office365
 - Google workspace
 - Netflix
 - ... (too numerous to list many)

Case Study: AWS

- **Virtualization technologies**
 - Xen
 - Amazon Nitro System
- **Specific virtualized resources**
 - EC2 instances (VMs)
 - Elastic Block Store (Storage)
 - Virtual Private Cloud/Virtual Private Networking
 - Managed Kubernetes: EKS
- **Impact**
 - Scalable on-demand computing
 - Hybrid Cloud Capabilities

Case Study: Azure

- **Virtualization technologies**
 - Hyper-V
- **Specific virtualized resources**
 - Azure Virtual Machines
 - Azure Disk Storage
 - Azure Virtual Networks
 - Managed Kubernetes: AKS
- **Impact**
 - Integration with the Microsoft ecosystem
 - Hybrid Cloud Capabilities

Case Study: Google GCP

- **Virtualization technologies**
 - KVM hypervisor
- **Specific virtualized resources**
 - Compute Engine (VMs)
 - Google Cloud Storage (Disks)
 - VPC, VPN (Networking)
- **Impact**
 - Integration with the Microsoft ecosystem
 - Hybrid Cloud Capabilities

Summary: Virt pros for cloud

- **Resource Optimization**
 - Improved utilization
 - Dynamic allocation
- **Scalability**
 - On-demand provisioning
 - Rapid deployment
- **Isolation**
 - Fault isolation
 - Enhanced security

Summary: Virt cons for cloud

- **Performance overhead**
 - Latency
 - Resource contention
 - When overcommitting
- **Security risks**
 - Hypervisor vulnerabilities
 - Isolation breakdowns
- **Licensing & Compliance**
 - Conditions imposed by some vendors make it complex
 - Regulatory compliance may require specific controls
 - Not available in the virtualized environment
- **Specialized knowledge**
 - Introduces its own complexity
 - Enhanced security


Mitigating challenges

- **Performance optimization**
 - Resource Monitoring
 - Hardware acceleration (VT-x/AMD-V)
 - When overcommitting
- **Security measures**
 - Security best practices
 - Frequent patching
 - Network segmentation
- **Licensing & Compliance**
 - Negotiations and compliance audits
 - Clear documentation
- **Specialized knowledge**
 - Automation tools
 - Training

Future trends

- **Edge computing**
 - Small devices close to the data
- **Serverless architectures**
 - Developer writes functions
 - Does not worry about servers/state
 - Need fast function activation and setup
- **Hyperconvergence**
 - Integration of compute, storage and networking into a single system
- **AI (of course)**
 - Motivation for innovations to efficiently support AI tasks.
 - E.g., GPU virtualization

Edge computing



Virtualization and Edge Computing

- **Processing data at the edge**
 - Small devices close to the data
- **Lightweight virtualization solutions**
 - Containers for small devices
 - MicroVMs
- **Use Cases**
 - IoT Devices
 - E.g. Autonomous vehicles, industrial sensors
 - Real time analytics
 - For critical applications, needing responsiveness

Serverless

- **Function as a Service (FaaS)**
 - Write and deploy code without caring how it is deployed
 - FaaS in charge of activating code and linking it to whatever services it needs
- **Event-driven execution**
 - Launch when certain events happen
 - E.g., arrival of API call
 - Dormant (not consuming resources) otherwise
- **Benefits**
 - Reduced operational complexity for developer
 - Cost efficiency
 - Automatic scaling

AI impact on virtualization

- **Need to optimize Virtualization for AI workloads**
- **GPU virtualization**
 - Same needs for sharing the resource
 - Same needs to keep an isolated environment
- **Intelligent resource management**
 - Optimize resource allocation
 - Cost efficiency
 - Automatic scaling

Best practices

- **Regular updates & maintenance**
 - Eliminate vulnerabilities
- **Resource monitoring and Optimization**
 - Prevent resource contention
 - Including the network
- **Security policies and procedures**
 - Access control, network segmentations, etc...
- **Disaster Recovery and Planning**
 - Backups, etc...

Key takeaways

- **No Cloud Computing without virtualization**
- **Different types of virtualization**
- **Need to understand pros and cons for proper decisions**
- **Technology is evolving**