



**kumori**  
SYSTEMS



**Writing Services:**  
The service model

# Goals

- Understand what the kumori service model entails
- Get proficiency on how to author kumori components and services
- Learn how to interact and deploy services through the axebow platform

<https://docs.kumori.systems/>



# Short glossary

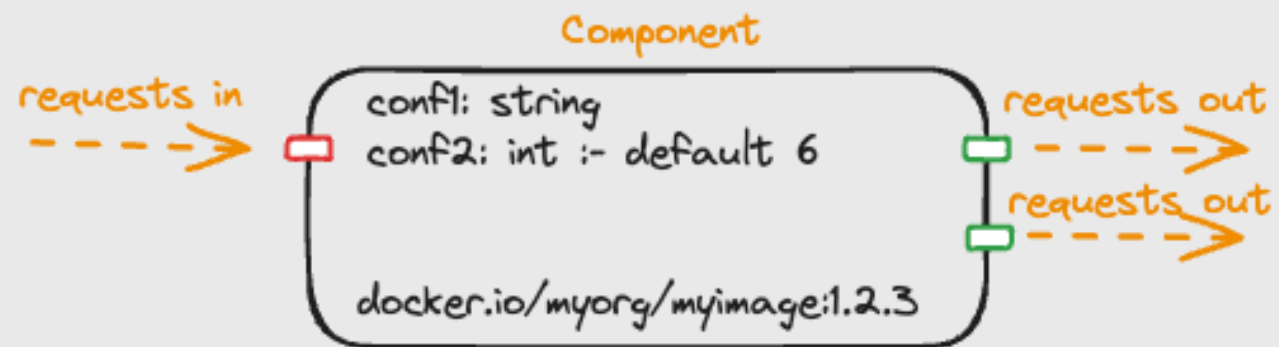
- Application
  - Program
  - Precise/unambiguous specification of how some system must behave
- Service
  - What one gets out of deploying, and running a program.
    - Holding state
    - Requiring life-cycle management
  - Typically through some endpoints
- Microservice
  - Service focused on delivering an atomic functionality
  - Typically part of a larger, more complex service.



# Elements of the Kumori Service Model

- ***Component***
  - Specifies the behavior of a microservice
    - Contains the code/application
    - In Kumori, it must be packaged as a docker image.
  - Specifies, among other things, how the component can be configured when activated
  - Also specifies how it communicates
    - Which dependencies it needs to communicate with
    - What endpoints it opens to accept requests from others





- Server channel: binding endpoint
- Client channel: connecting endpoint

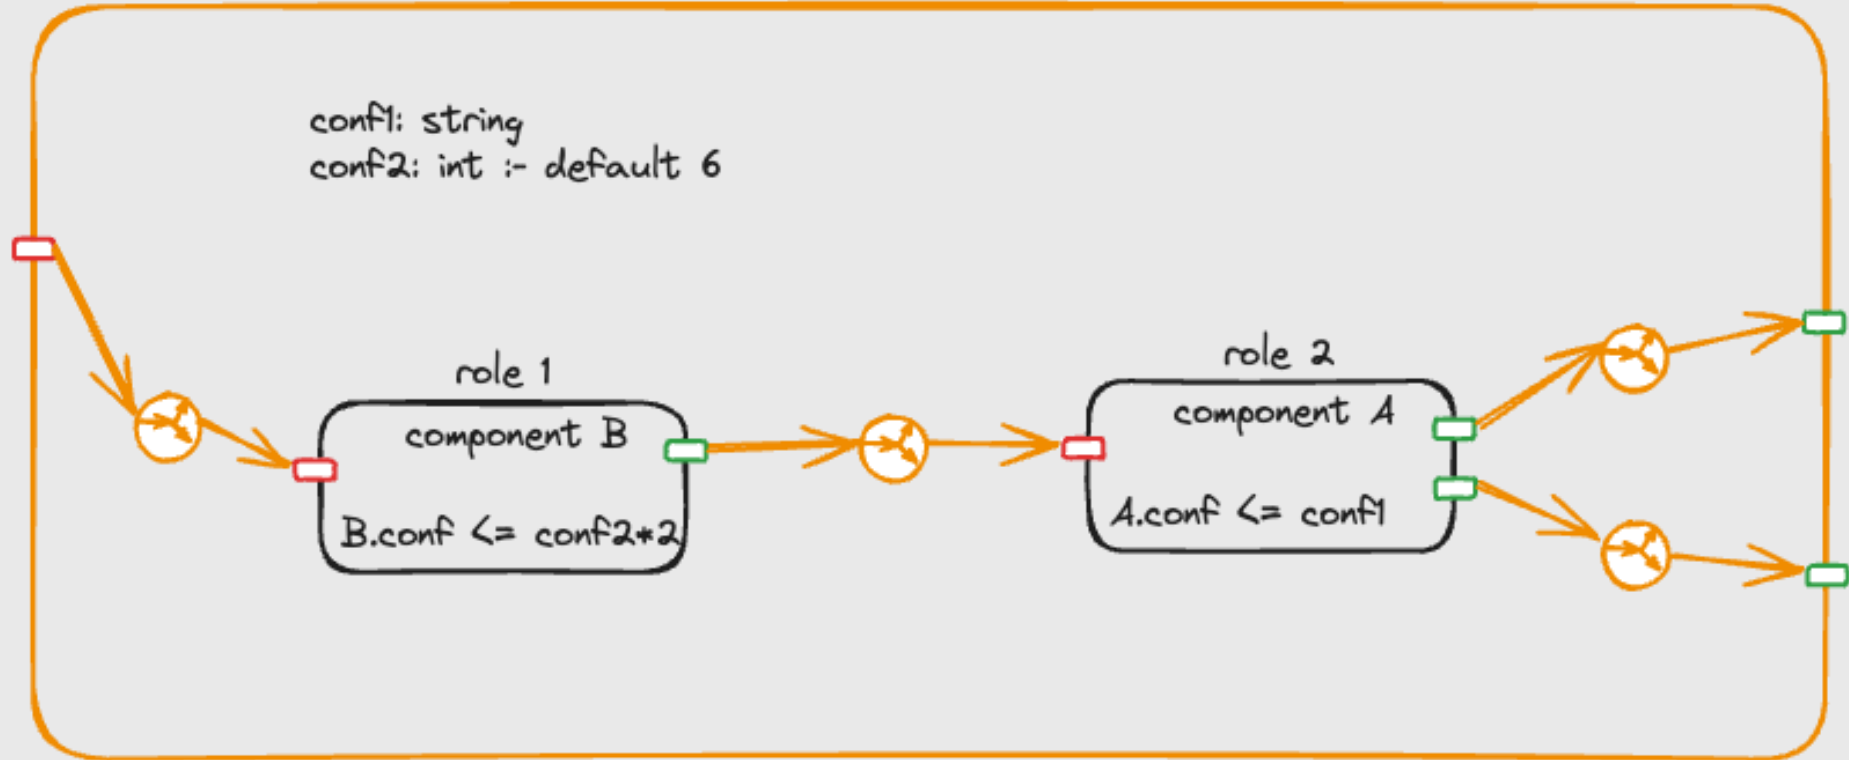


# Elements of the Kumori Service Model

- ***Service application***
  - Specifies how to structure multiple microservices within a complete service
    - Roles encapsulating components as microservices
    - Communication patterns between roles
      - Taking into account the component's endpoints
  - Specifies how the whole service is configured
    - With rules distributing that configuration to each one of the service's role
    - How each role distributes/transforms configuration for its component
  - Also specifies how it communicates
    - Which dependencies it needs to communicate with
    - What endpoints it opens to accept requests from others



## Service application



 Load balancer connector

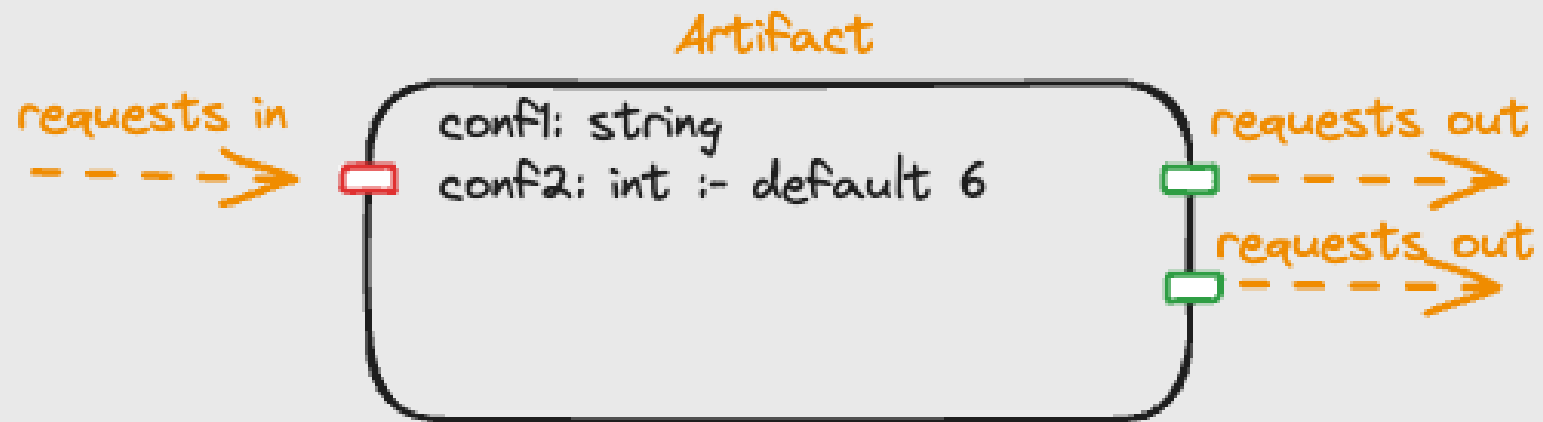




# Elements of the Kumori Service Model

- *Artifact:*
  - *Component*
  - *Service Application*
- *Common elements of artifacts*
  - Configuration
  - Communication endpoints (aka “channels”)







-  Server channel: binding endpoint
-  Client channel: connecting endpoint



# Elements of the Kumori Service Model

- *Deployment*
  - Sets the initial goal state of a service
    - Identifies the service application to be deployed
    - Provides concrete values for the configuration of the service
      - Following the configuration spec of the service application being deployed



# Specifying Kumori artifacts

- Uses a Domain Specific Language
- In v. 1.0: Based on cuelang
  - Internal DSL on top of CUE
- In v 2.0
  - External DSL, with VSCode support
  - Better tooling
- We will look at the v1.0 language
  - V2.0 still in the oven 😞



# Manifests

- An artifact is specified with a manifest for the artifact
- A manifest has the following general structure

```
{  
  spec: #Version  
  ref: #ManifestRef  
  description: _  
}
```

- Specification language is CUE
  - <https://cuelang.org>
- Manifests form an internal DSL within CUE



# Manifests

- *Artifact manifests are shipped within kumori modules*
- *Kumori module*
  - *Unit of versioning and distribution*
  - *A Kumori module has its own reference*
  - *A Kumori module is versioned*
- *A kumori module can contain many different manifests of artifacts*
  - *Each one within its own directory*
  - *Names of artifacts follow the path of the directory they are within*
  - *All artifacts within a module inherit the module's version, domain and module name (oc)*



# Manifests

```
{  
  spec: #Version  
  ref: #ManifestRef  
  description: _  
}
```

- The ***spec*** field specifies the version of the manifest specification
  - Determines the format for the rest of the manifest
- ***description*** contains the specific info needed for each kind of artifact
- The ***ref*** field uniquely identifies the artifact being defined.



# Manifests: the Ref field

```
#ManifestRef: {  
  version: *[1, 0, 0] | #Version  
  domain: string  
  module: string  
  name: string  
  kind: string  
}
```

- Field ***version*** follows semantic versioning
- Field ***kind*** specifies the artifact kind (***component***, ***service***)
- Field ***domain*** identifies the owner of the artifact (as a domain name)
- Field ***module*** identifies the module that distributes the artifact
- Field ***name*** completes the ref ID of the artifact.



# Publishing/reproducibility

- Once a module has been published, it does not change
  - Module immutability: all artifacts within a module's version stay constant
- The *ref* uniquely identifies an artifact
  - Its name derived from the full module's name
- A module can be ported
  - ... and so all its artifacts
- A set of modules can be bundled
  - To move it or to prepare it for deployment





# Publishing/reproducibility

- It is not necessary to have a local registry of modules
  - But there is a caching mechanism
  - Local modules can be used
- Modules are registered currently as npm packages
  - A CLI tool exists to facilitate this.
  - The npm scope must coincide with the domain of the module.
    - Thus, the publisher must have credentials to publish under the scope



# Tooling (advance)

- Tool: *kam*
- Installation:  

```
npm i -g @kumori/kam
```
- Usage:
  - **kam** mod init domain/modname
  - ...
- Can be used to run *cue*
  - **kam** cue ...
- Needs nodejs



# Preliminaries: the CUE language

- Configuration specification language
  - Superset of JSON
    - Any JSON is CUE
    - Any YAML can be converted to CUE
- Based on logic programming
- Lets you define data schemas
  - With more generality than JSON Schema or OpenAPI
  - Thus factoring out repetitive patterns
- Has sufficient expressiveness to derive values out of other values



# CUE

- Implements a set type system
  - Thus, individual elements within a type are also types
  - CUE does not distinguish between them (except for “exporting”)
- “type” specifications are simply restrictions imposed on a field

## Schema:

```
ciudad: {  
  nombre: string  
  pop: int  
  capital: bool  
}
```

## Specialization:

```
cGrande: {  
  nombre: string  
  pop: >5  
  capital: true  
}
```

## Data:

```
moscu: {  
  nombre: "Moscó"  
  pop: 11.92  
  capital: true  
}
```



# CUE: basic types

## Basic types

int

number

string

bool

null

## Special

—

any/top

—|—

bottom/  
nothing

## Structured

[...]

Any array

{...}

Any struct



# CUE

- The same field can be specified multiple times
  - Each time adds a restriction to the contents the field can have
  - If the set of restrictions on a field end up being incompatible, the field equals `_|_`
    - ...which makes the structure it is part of also be `_|_`
- Order of field specifications is unimportant

a: int  
a: >0

a: int  
a: > 0  
a: 5

a: < 100

a: 6  
Conflict



# CUE: operations

- *Disjunction* (or union...)

- Given two types, The resulting type is the union of both
- Operator |
- Commutative and associative

```
ob1: "something" | "another"
```

```
ob2: string | uint
```

```
ob3: ob1 | ob2
```

```
ob3: 56
```



# CUE: unification

- Given two types, A, B
  - Find the “largest” type, C, such that C is part of A and C is part of B.
- “largest”?
  - Let C be included in A and B,
    - C is the largest iff
      - For any D included in A and B, D is included in C.
- Operator. &.

```
ob1: {  
  a: int  
}
```

```
ob2: {  
  b: string  
}
```

```
ob3: ob1 & ob2
```

```
ob3: {  
  a: int  
  b: string  
}
```





# CUE: unification

- Two restrictions on the same field function as a unification operation
  - Commutative and associative operation

```
ob1: {  
  a: int  
}
```

```
ob2: {  
  b: string  
}
```

```
ob3: ob1  
ob3: ob2
```

```
ob3: {  
  a: int  
  b: string  
}
```



# CUE: unification

```
ob1: {  
  a: string  
  b: 3  
}
```

```
ob2: {  
  a: "something"  
}
```

```
ob3: ob1  
ob3: ob2
```

```
ob3: {  
  a: "something"  
  b: 3  
}
```



# CUE: defaults

```
ob1: {  
  a: string  
  b: int | *3  
}
```

```
ob2: {  
  a: "something"  
}
```

```
ob3: ob1  
ob3: ob2
```

```
ob3: {  
  a: "something"  
  b: int | * 3  
}
```



# CUE: defaults

```
ob1: {  
  a: string  
  b: int | *3  
}
```

```
ob2: {  
  a: "something"  
  b: 67  
}
```

```
ob3: ob1  
ob3: ob2
```

```
ob3: {  
  a: "something"  
  b: 67  
}
```



# CUE: definitions

- Another way to specify fields
- A field is a definition if its name starts with “#”
- A definition introduces a closed structure

## Closed Structure

- Cannot be unified with anything that would introduce new fields
  - Si se unifica con otra estructura, la segunda no puede introducir un campo nuevo



# CUE: estructuras cerradas

```
#A: {  
  pop: int  
  name: string  
}
```

```
B: {  
  capital: bool  
}
```

madrid: #A & B

CONFLICT



# CUE: subsumptions/mezclas

```
#A: {  
  pop: int  
  name: string  
}
```

```
B: {  
  capital: bool  
}
```

```
madrid: {  
  #A  
  B  
}
```

```
madrid: {  
  pop: int  
  name: string  
  capital: bool  
}
```



# CUE: export

```
#A: {  
  pop: int  
  name: string  
}
```

```
B: {  
  capital: bool  
}
```

```
madrid: {  
  #A  
  B  
} & {  
  pop: 3  
  name: "Madrid"  
  capital: true  
}
```

cue export

ERROR





# CUE: export

```
#A: {  
  pop: int  
  name: string  
}
```

```
#B: {  
  capital: bool  
}
```

```
madrid: {  
  #A  
  #B  
} & {  
  pop: 3  
  name: "Madrid"  
  capital: true  
}
```

cue export

OK

```
madrid: {  
  pop: 3  
  name: "Madrid"  
  capital: true  
}
```



# CUE: referencias

```
#A: {  
  spec: #B  
  alt: spec.capital  
}
```

```
spain: #A & {  
  spec: {  
    capital: "Madrid"  
  }  
}
```

```
#B: {  
  capital: string  
}
```

```
spain: {  
  spec: {  
    capital: "Madrid"  
  }  
  alt: "Madrid"  
}
```



# Kumori Modules

- CUE code organized within modules
- Modules are the distribution units
- MODULE == special structure folder
  - Subfolder ***cue.mod***
    - Managed by kumori tool, needed for CUE loader: DO NOT TOUCH
  - File ***kmodule.cue*** contains fields (see next slide),
    - *Relevant names (domain, module, version)*
    - Dependencies
      - With query/lock
      - With an alias that can be used in imports
    - Checksums
  - Common files for components and services
    - Auto-generated: DO NOT TOUCH



cue.mod

inbound

.gitignore

.gitlab-ci.yml

README.adoc

componentref.cue

kmodule.cue

service\_artifact.cue

serviceref.cue

utils.cue

```
package kmodule

{
    local: true
    domain: "kumori.systems"
    module: "builtins/inbound"
    cue: "v0.4.3"
    version: [
        1,
        3,
        1,
    ]
    dependencies: "kumori.systems/kumori": {
        target: "kumori.systems/kumori/@1.1.6"
        query: "1.1.6"
    }
    sums: "kumori.systems/kumori/@1.1.6": "jsXEYdYtlen2UgwDYbUCGWULqQIigC6HmkexXkyp/Mo="
    spec: [
        1,
        0,
    ]
}
```



# Kumori packages and manifests

- Kumori manifests reside within a kumori module (thus a CUE module)
- Given a KMODULE, several artifacts can be within
- The manifests for component artifacts reside within a “component” CUE package
- The manifests for service artifacts reside within a “service” CUE package
- Given the CUE loader
  - It is possible to share parts of manifests among multiple artifacts of the same kind



cue.mod

inbound

.gitignore

.gitlab-ci.yml

README.adoc

componentref.cue

kmodule.cue

service\_artifact.cue

serviceref.cue

utils.cue

```
// Automatically generated file. Do not edit.  
package component
```

```
import (  
  k "kumori.systems/kumori:kumori"  
  m "...:kmodule"  
)
```

```
#Artifact: k.#Artifact & {  
  spec: m.spec  
  ref: {  
    local: m.local  
    version: m.version  
    if m.prerelease != _|_ {  
      prerelease: m.prerelease  
    }  
    if m.buildmetadata != _|_ {  
      buildmetadata: m.buildmetadata  
    }  
    domain: m.domain  
    module: m.module  
    kind: "component"  
  }  
}
```



# Packages

- Modules are structured within packages
  - Three package names used within kumori modules
    - *kmodule*
    - *component*
    - *service*
  - Also, the main “*kumori*” package to distribute the model restrictions
    - Only used by kumori code
- Package loader works following folder structure.
  - We take advantage of it.



cue.mod

inbound

.gitignore

.gitlab-ci.yml

README.adoc

componentref.cue

kmodule.cue

service\_artifact.cue

serviceref.cue

utils.cue

```
// Automatically generated file. Do not edit.  
package service
```

```
import (  
  k "kumori.systems/kumori:kumori"  
  m "...:kmodule"  
)
```

```
#Artifact: k.#Artifact & {  
  spec: m.spec  
  ref: {  
    local: m.local  
    version: m.version  
    if m.prerelease != _|_ {  
      prerelease: m.prerelease  
    }  
    if m.buildmetadata != _|_ {  
      buildmetadata: m.buildmetadata  
    }  
    domain: m.domain  
    module: m.module  
    kind: "service"  
  }  
}
```

**kumori**  
SYSTEMS





cue.mod

inbound

.gitignore

.gitlab-ci.yml

README.adoc

componentref.cue

kmodule.cue

service\_artifact.cue

serviceref.cue

utils.cue

Name

..

placeholder.cue

```
package service

import (
    k "kumori.systems/kumori"
)

// Definition of a list of IPs (regular IPs or CIDRs, defined via regular expression)
let ip = "(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\\.?(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\\.?(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)\\.?(25[0-5]|2[0-4][0-9]|[01]?[0-9][0-9]?)$"
let cidr = "\\(ip)/([0-9]|[1-2][0-9]|3[0-2])"
#ipcidr: =~ "^\\(ip)\\(cidr)$"
#IPList: [...(#ipcidr)]

#Artifact: {
    ref: name: "inbound"

    description: {
        builtin: true
        config: {
            parameter: {
                type: *"https" | "tcp"
                if type == "https" {
                    clientcert: bool | *false
                    if clientcert == true {
                        certrequired: bool | *false
                    }
                }
                websocket: bool | *false
                // Name of the header (for example, "X-Real-IP") where the remote address
                // will be set. This may not be the physical remote address of the
                // peer if the address has been inferred from the x-forwarded-for
                // (depends on cluster configuration)
                remoteaddressheader: string | *""
                // Clean the x-forwarded-for header
                cleanxforwardedfor: bool | * false
                // An http inbound can declare a list of allowed or denied IPs
                //
                // This syntax is preferable, because it prevents the simultaneous
                // existence of "allowediplist" and "deniediplist", but an error occurs
                // during the service spread.
                //   accesspolicies?: { allowediplist: #IPList } | { deniediplist: #IPList }
                //
                // (currently not supported for tcp inbounds)
                //
                accesspolicies?: {
                    allowediplist?: #IPList
                    deniediplist?: #IPList
                }
            }
        }
    }
    resource: {
        if config.parameter.type == "https" {
            servercert : k.#Certificate
            serverdomain : k.#Domain
            clientcertca?: k.#CA
        }
        if config.parameter.type == "tcp" {
            port: k.#Port
        }
    }
    srv: client: inbound: _
}
```

kumori  
SYSTEMS

# Importing packages

- Within a CUE file it is possible to import other packages
  - Typically from other kumori modules holding artifacts
  - Potentially from the kumori module holding the model spec itself.
- Syntax similar to GOlang

```
import (  
  k    "kumori.systems/kumori/kmv3_0_1"  
      "kumori.systems/integrations/hazelcast/v3/server"  
  mngmnt "kumori.systems/integrations/hazelcast/v3/management_center:component"  
  mc    "mod.local/subdir:component"  
)
```

- Can provide alias (imports 1, 3 & 4)
  - Implicitly, the basename of the import (import 2: server)
  - Must provide the kind of artifact (3)
  - Can refer to artifacts within the current module as “mod.local”
- A package fields are available qualifying them with the alias of the import.



# Components

- Description with several sections:
  - Microservice
  - Configuration
  - Size
  - Code

```
description: {  
  srv: {...}  
  config: {...}  
  size: {...}  
  code: {...}  
}
```



# Component: microservice

- Channel description:
  - **Client** channel (dependencies)
  - **Server** channel
    - Must specify the port
  - **Duplex**
    - Sort of shortcut representing a client and a server
    - Exposes the endpoints linked to it (IP:port)
    - Complex protocols
      - E.g., DDBB, distributed runtimes (BEAM,...)
- All names are local to the component
  - Independent of the service where the component is used



# Component: microservice

```
srv: {  
  client: [string]: #Channel  
  server: [string]: #Server  
  duplex: [string]: #Server  
}  
  
#Channel: {  
  protocol: "udp" | "tcp" | *"http" | "grpc"  
}  
  
#Server: {  
  #Channel  
  port: uint16 | *80  
}
```



# Component: service discovery

- What is?
  - Process by means of which the code discovers at run time how to contact a dependency.
- KPaaS uses DNS mechanisms
  - Uses domain names → names of client channels
    - Local to the component.
      - Known at “compile time”
    - KPaaS injects dependencies at deployment time
    - KPaaS modifies info as it changes
- In KPaaS dependencies are mainly client channels
  - Also duplex channels, though they work differently



# Component: basic service discovery

- KPaaS solves dependencies as follows
  - Given a client channel, **C**,
    - nslookup O.C → IP(s) of servers
    - nslookup -S O.C → IP:port of servers
    - Tag “0” refers to the first target of a dependency
      - A channel can point to several endpoints, each one has a different tag number
      - /kumori/config.json contains information about tags
- Bonus
  - Reserved channel: **self**
  - nslookup self → IP de la instancia
    - Useful for binding to a specific external interface for the container (**bind**)
- Bonus II
  - It is also possible to resolve channel “C”
    - The result is address 127.0.0.N (We will see this later)



# Component: Configuration

- Under field config
- Two subsections: parameter, resource
- parameter is a dictionary/tree
  - Arbitrary JSON
- resource is also a dictionary/tree
  - Leaves satisfy #Resource

config: parameter: [string]: \_





# Componente: Configuración, recursos

config: resource: #ResourceTree

#Resource: #Volume | #Secret  
#Volume: #Persistent | #Volatile

#Secret: {secret: string}  
#Volatile: {size: uint, unit: string}  
#Persistent: {volume: string}



# Component: Size

- Simple declaration of the “vertical” resources that instances of a component need

```
size: #ComponentSize
```

```
#ComponentSize : {  
  mincpu:      #ResourceAmount & {kind: "cpu"}  
  minbandwidth?: (number & >= 0 & <= bandwidth.size)  
  bandwidth:   #ResourceAmount & {kind: "bandwidth"}  
}
```

```
#ResourceAmount {  
  kind:      #UnitKinds  
  size:      number & >= 0  
  unit:      #Units[kind]  
}
```

```
#Units: {  
  storage: "k" | "M" | "G" | "T" | "P" | "E" | "Ki" | "Mi" | "Gi" | "Ti" | "Pi" | "Ei"  
  cpu: "m" | ""  
  ram: "G" | "M" | "Gi" | "Mi"  
  bandwidth: "G" | "M" | "Gi" | "Mi"  
}
```



# Component: Code

- Description of how the containers should be activated
- As a dictionary
  - The keys are the container names
  - Values describe how each container should be configured
    - Its image
    - How to map configuration to environment and filesystem entities

```
code: [nm=string]: #Container & {name: nm}
```



# Component: Code

```
#Container: {  
  name: string  
  image: #Image  
  entrypoint?: [...string]  
  cmd?: [...string]  
  user?: {  
    userid: uint16  
    groupid: uint16  
  }  
  mapping: #Mapping  
}
```

```
#Image: {  
  hub: #Hub | * {  
    name: "registry.hub.docker.com"  
    secret: ""  
  }  
  tag: string  
  digest?: string  
}  
  
#Hub: {  
  name: string  
  secret: string  
}
```



# Component: Mappings

- Two types: Environment variables, or file system
  - Use cue references from the config section.
- Mechanism to convert config to constructs usable by the code within the container

```
#Mapping: {  
  filesystem: [...#FileMap | #FolderMap]  
  env: #EnvMap  
}
```

```
#EnvMap: {  
  [string]: {value: string} | #Secret  
}
```



# Component: Mappings

- Two filesystem mapping types: Folders and Files
- Lets developers define arbitrary folder trees
  - A folder must map to a tree or to a Volume resource (by resource name)

```
#FileMap: {  
  path: string  
  mode: uint16 | *0o644  
  data: _//jsonable  
  format: *"text" | "json" | "yaml"  
} | {  
  path: string  
  mode: uint16 | *0o644  
  #Secret  
}  
#FolderMap: {  
  path: string  
  { {tree: [...#FileMap | #FolderMap]} | #Volume }  
}
```



# Component: Spread & Example

Frontend de calculator cache

<https://gitlab.com/kumori-systems/community/examples/calc-cache>

Worker de calculator-cache



# Service

- A service's description has 5 sections
  - Microservice interface
  - Configuration
  - Roles
  - Connectors
  - Links
- NOTE: Components and Artifacts share the two first sections

```
description: {  
  srv:    {...}  
  config: {...}  
  role:   {...}  
  vset:   {...}  
  connector: {...}  
  link:   {...}  
}
```





# Service: Roles

- A **role** is a way of specifying how to deploy a component within a service application.
- The roles section within a service app is structured as a dictionary whose names are the role names.

```
role: [rn=string]: #Role & {name: rn}
```

```
#Role: {  
  artifact: #Artifact  
  name:    string  
  meta:    {...}  
  config: #Configurable  
}
```

```
#Configurable: {  
  parameter: [string]: _  
  resource:  [string]: #Resource  
  resilience: uint | *0  
  scale: #ScaleSpec  
}
```



# Service: Topology

- Contains connectors and links
- A connector allow linking from client to server channels within a service's role
  - Can be visualized as two endpoints:
    - Client endpoint: role client channels attach to it
    - Server endpoint: attaches to role server channels
- Connector types
  - Two types
    - LB: Load Balancer (rr)
    - FULL: complete connector/Full
    - Determine how to deal with the multiplicity of instances in the roles they attach to
      - Condition how service discovery ends up working



# Services: Links

- They are the edges of the topology
- Only the following attachments are allowed
  - From a role client channel to a connector client endpoint
  - From a connector server endpoint to a role server channel
    - When the connector is full and absend links to the client enpoint
      - also allow link to a duplex server channel
- ***self*** is special:
  - From one of the service's server channels to a conector's client endpoint
  - From a connector's server endpoint to one of the services' client endpoints
  - Only attaches to LB connectors



# Service: Topology

```
connect: {  
  as: "lb" | "full"  
  from: [rn=#roles | "self"]: role[rn].srv.#clients  
  to: [rn=#roles]: role[rn].#servers: {  
    meta: {...}  
  }  
}
```



# Service: Connectors and discovery

- LB Connector
  - Represented by a balanced IP
  - Channels connected to the client endpoint of an LB resolve the balanced IP
    - Its port is always 80
- FULL Connector
  - Channels connected to the client endpoint of the connector resolve all IP's of the server channels to which the server endpoint connects.
  - When resolving the SRV records, all IP:port are returned
  - When only the server endpoint is connected, it must connect to a duplex channel
    - Resolving the A record of the duplex channel returns all Ips of the instances
    - Resolving the SRV record of the duplex channel returns all IP:port of the instances



# Configuration Spread

- Same as for components: using CUE references and *unification*
  - Within the role record
    - At the config field: left hand side is the config for the role's artifact, right hand side is a reference to the service app fields
- Referencing, mainly, the configuration section of the service.

```
role: myrole: rsize: {  
  resilience: config.parameter.maxfailures  
}
```

```
role: myrole: artifact: myComponentManifest  
role: myrole: cfg: {  
  parameter: color: config.parameter.background  
}
```



# Deployment

- A deployment manifest provides
  - A reference to an artifact to be deployed
  - A set of concrete values for the configuration parameters of the artifact being deployed
    - Including data informing about the number of instances to be deployed for each role
- With a deployment manifest
  - A service initial revision can be deployed
  - A new revision can update a previously running revision

```
#Deployment: {  
  artifact: #Artifact  
  meta : {...}  
  config: {...}  
}
```

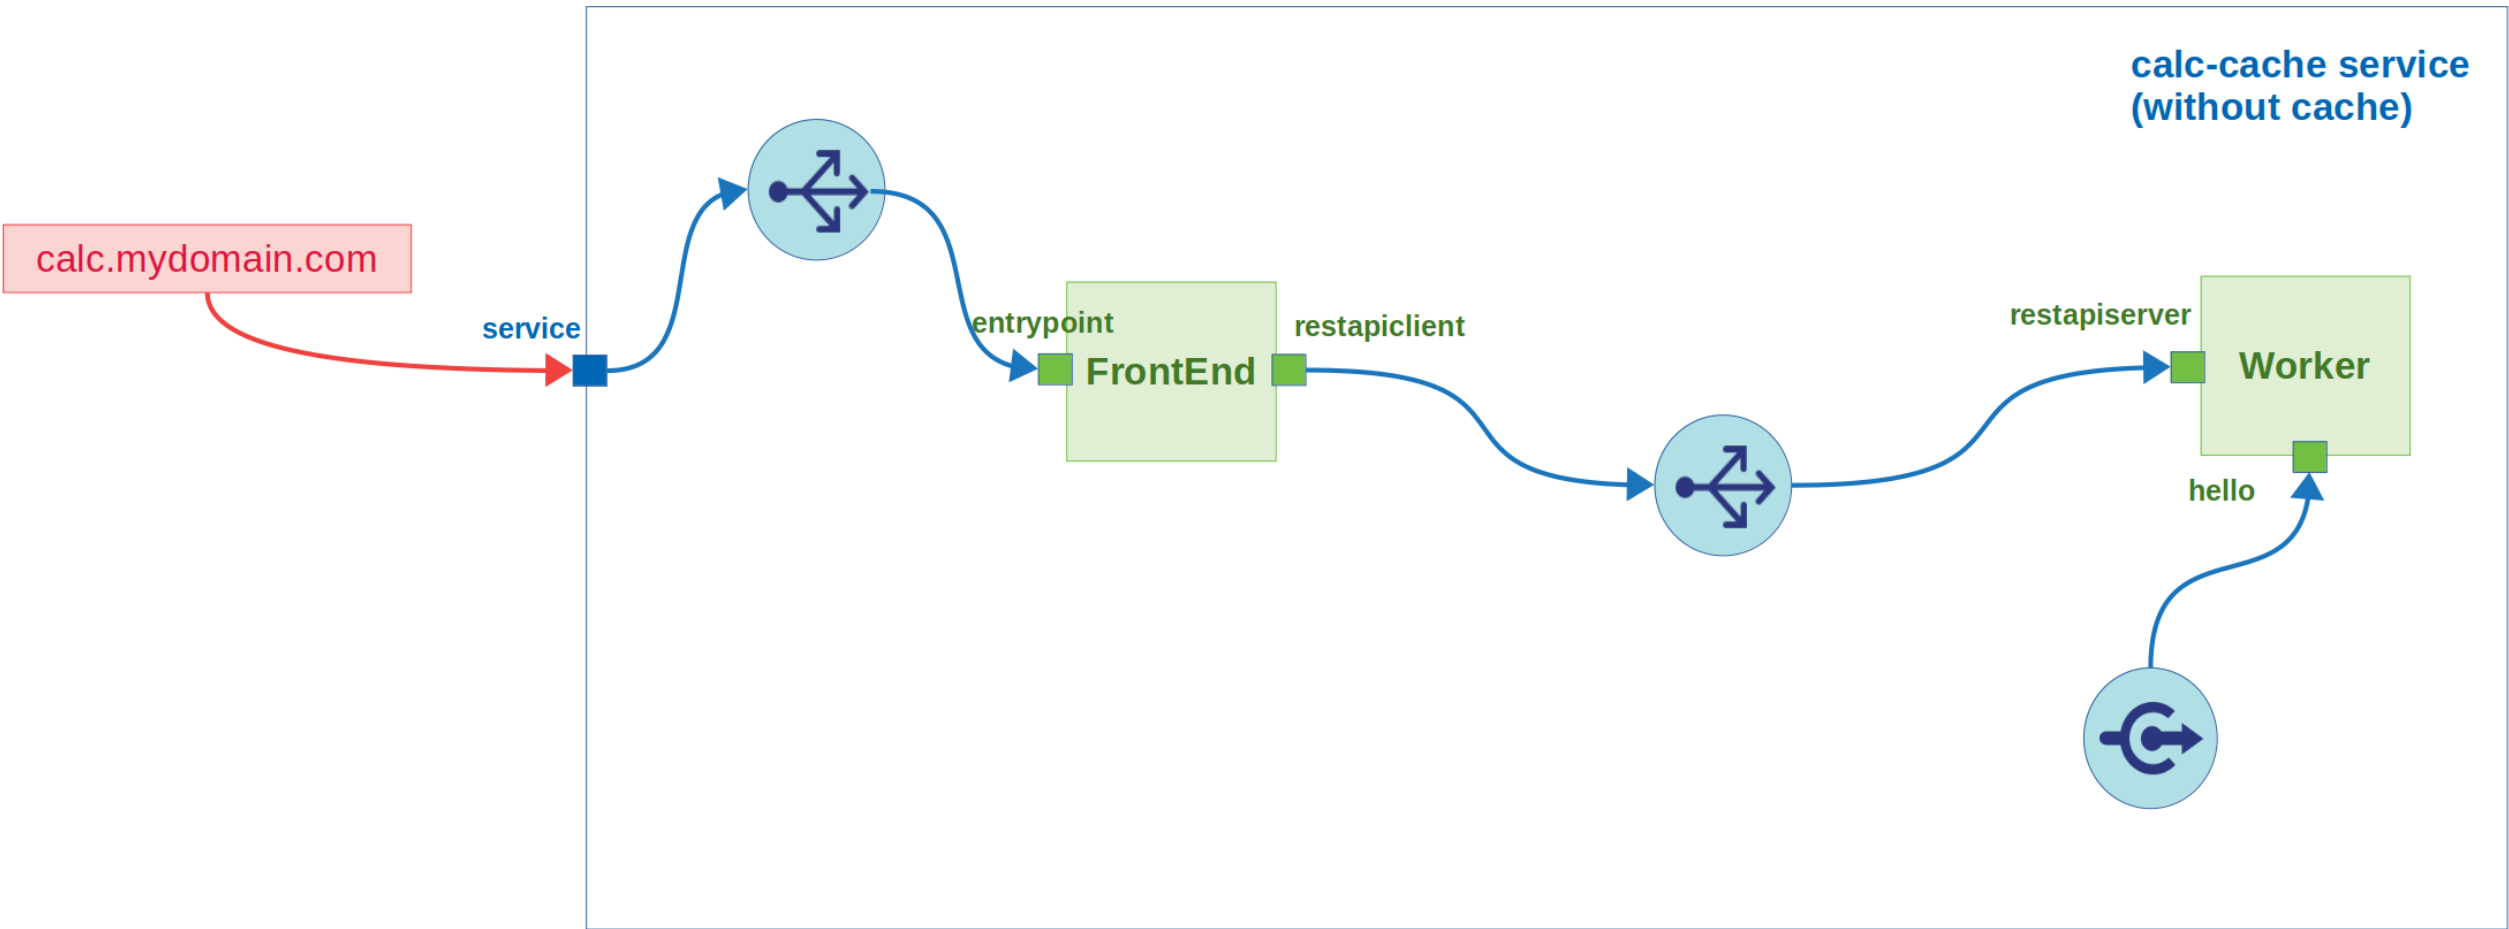


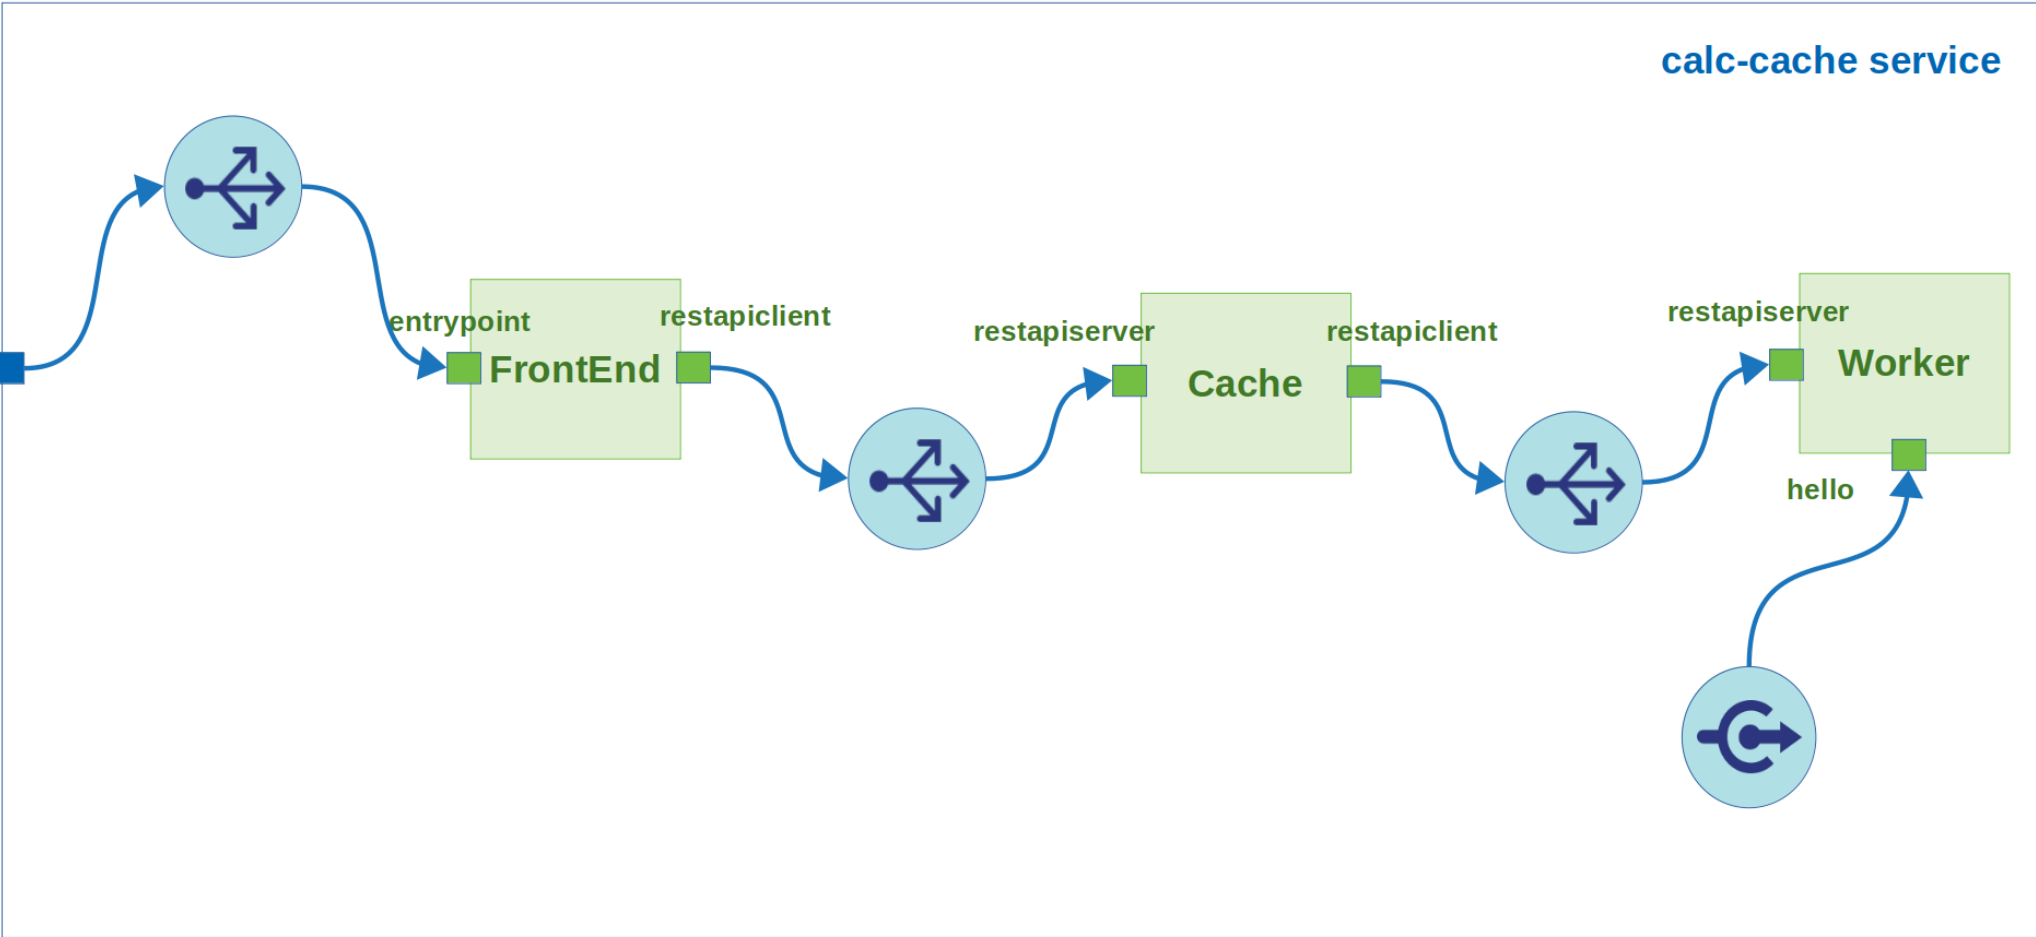
# Example: calculator-cache

- Toy example
- A simple calculator but with a possible twist
  - Frontend accepts requests and passes them to a worker
    - Worker has a dúplex cannal but this toy example does not need it.
  - A variation has the frontend request a value from the cache instead
    - The cache will contact the worker if need be
- A total of 2 roles in the simple case
  - 3 in the “more complex” case.









Turn to editor



# kam

- Setting up the working environment

- Create the working directory
  - mkdir tutorial
  - cd tutorial

kam mod init kumori.examples/calccache

- Creates a set of files

- Sets it up with some defaults

kam mod dependency kumori.systems/kumori

```
componentref.cue
cue.mod
kmodule.cue
serviceref.cue
```

```
{
  domain: "kumori.examples"
  module: "calccache"
  cue: "0.4.2"
  version: [
    0,
    0,
    1,
  ]
  dependencies: {}
  sums: {}
  spec: [
    1,
    0,
  ]
}
```



# kam

- Setting up the working environment

- Create the working directory
  - mkdir tutorial
  - cd tutorial

```
kam mod init kumori.examples/calccache
```

- Creates a set of files

- Sets it up with some defaults
- Then:

```
kam mod dependency kumori.systems/kumori
```

```
componentref.cue  
cue.mod  
kmodule.cue  
serviceref.cue
```

```
{  
  domain: "kumori.examples"  
  module: "calccache"  
  cue: "0.4.2"  
  version: [  
    0,  
    0,  
    1,  
  ]  
  dependencies: {}  
  sums: {}  
  spec: [  
    1,  
    0,  
  ]  
}
```



# kam

- Setting up the working environment

- Create the working directory
  - mkdir tutorial
  - cd tutorial

kam mod init kumori.examples/calccache


- Creates a set of files

- Sets it up with some defaults
- Then:

kam mod dependency kumori.systems/kumori

```
componentref.cue
cue.mod
kmodule.cue
serviceref.cue
```

```
{
  domain: "kumori.examples"
  module: "calccache"
  cue: "0.4.2"
  version: [
    0,
    0,
    1,
  ]
  dependencies: {}
  sums: {}
  spec: [
    1,
    0,
  ]
}
```



```
{
  domain: "kumori.examples"
  module: "calccache"
  cue: "0.4.2"
  version: [
    0,
    0,
    1,
  ]
  dependencies: "kumori.systems/kumori": {
    target: "kumori.systems/kumori/@1.1.7"
    query: "latest"
  }
  sums: "kumori.systems/kumori/@1.1.7": "kPdupjoBs/7ZLsDSsJCXEoY4Su+L3LCpbXMK4nBwbQY="
  spec: [
    1,
    0,
  ]
}
```





**kumori**  
SYSTEMS



The next  
cloud platform